

Name: Nikhil Kumar Gattu

Mail: gattun@oregonstate.edu

ONID: 934615235

1. Tell what machine you ran this on

Ans: I ran the experiment on flip4.engr.oregonstate.edu server, which supports OpenMP version 201511.

2. What performance results did you get?

Ans:

```
...4:~ — ssh gattun@access.engr.oregonstate.edu
[flip4 ~ 78$ bash loop.bash
OpenMP version 201511 is supported here
For 1 threads, Peak Performance = 160.92 MegaMults/Sec
OpenMP version 201511 is supported here
For 4 threads, Peak Performance = 615.54 MegaMults/Sec
flip4 ~ 79$ █
```

1 thread: 160.92 MegaMults/Sec
4 thread: 615.54 MegaMults/Sec

3. What was your 1-thread-to-4-thread speedup?

Ans: $S = 615.54 / 160.92 = 3.82$

4. Your 1-thread-to-4-thread speedup should be less than 4.0. Why do you think it is this way?

Ans: I think real-world performance is always affected by:

- Overhead in thread management (creating and scheduling threads)
- Memory bandwidth limitations
- CPU cache sharing
- Small portions of the program that remain sequential (as explained by Amdahl's Law)

These factors prevent us from reaching perfect linear scaling.

5. What was your Parallel Fraction, F_p ? (Hint: it should be less than 1.0, but not much less.)

Ans: $F_p = 4/3 \times (1 - 1/3.82)$

$= 4/3 \times (1 - 0.262)$

$= 4/3 \times 0.738$

$= 0.984$

This means about 98.4% of the code is parallelizable, which is excellent and expected since the array multiplication is a highly parallel task.