

Name: Nikhil Kumar Gattu

Mail: [gattun@oregonstate.edu](mailto:gattun@oregonstate.edu)

ONID: 934615235

1. What machine you ran this on


➔ I ran on Flip3 server in my personal laptop (MacOS).

2. Show the 2 tables of performances for each array size and the corresponding speedups.

Table 1: Array Multiplication Performance and Speedup			
Array Size	Non-SIMD Perf (MegaMul/s)	SIMD Perf (MegaMults/s)	Speedup (SIMD / Non-SIMD)
1000	360.74	1779.19	4.93
6000	359.19	2045.55	5.69
50000	355.48	1885.44	5.3
80000	356.16	1872.87	5.26
160000	355.15	1748.21	4.92
550000	353.59	1748.9	4.95
1000000	350.65	1683.66	4.8
2000000	344.47	1130.28	3.28
6000000	340.02	845.17	2.49
8000000	341.35	844.31	2.47

Table 2: Array Multiply + Reduction Performance and Speedup			
Array Size	Non-SIMD Perf (MegaMul/s)	SIMD Perf (MegaMults/s)	Speedup (SIMD / Non-SIMD)
1000	362.44	1798.56	4.96
6000	364.8	2060.6	5.65
50000	365.36	1959.85	5.36
80000	365.36	1988.78	5.44
160000	364.95	1918.72	5.26
550000	363.55	1860.21	5.12
1000000	363.1	1746.76	4.81
2000000	359.03	1619.59	4.51
6000000	354.68	1103.09	3.11
8000000	356.31	1082.68	3.04

3. Show the graphs (or graph) of SIMD/non-SIMD speedup versus array size (either one graph with two curves, or two graphs each with one curve)



#### 4. What patterns are you seeing in the speedups?

- ➔ In both results, array multiplication and array multiplication with reduction the speedup from using SIMD is initially high, peaking around 5 to 5.7x. This trend is strongest with small to medium array sizes (from 1K to ~500K). As the array size increases beyond that point, the speedup begins to gradually decline. By the time we reach 8 million elements, speedup drops to around 2.4 to 3.0x, depending on the operation.

#### 5. Are they consistent across a variety of array sizes?

- ➔ The speedups are relatively consistent within each size range, especially for small and medium arrays. There's a smooth, predictable curve, the speedup starts high and gradually flattens out or declines. This shows that SIMD performance benefit is consistent within resource-friendly sizes, but the trend changes as we scale up.

#### 6. Why or why not, do you think?

- ➔ This change in pattern likely stems from memory hierarchy and bandwidth limitations. For small arrays, most data fits in CPU cache, so SIMD instructions benefit from low-latency memory access. As the array size grows, data must be loaded from main memory, increasing latency and reducing the effectiveness of vectorization. Additionally, I think the overhead of loop control, cache misses, and lack of memory parallelism limits the gain from SIMD, even though it's still faster than non-SIMD execution.